

Regression in PySpark on Databricks platform

- This notebook covers developing a machine learning model in **PySpark** framework on **Databricks**.
- Dataset used for analysis is Sarah Gets a Diamond (<http://store.darden.virginia.edu/sarah-gets-a-diamond>) taken from University of Virginia, Darden Business Publishing. I have purposefully and randomly removed values from 30 rows for two different columns, so that different data transformation techniques in PySpark framework can be applied.
- A regression model is developed on this dataset to calculate the price(Y variable) of a diamond based on its multiple physical attributes/features(X variables).
- Different data transformations, data imputation techniques, data aggregation in PySpark environment are applied for cleaning, imputing, pre-processing, and extracting summary statistics on model development data.
- Regression model development and tuning is performed under PySpark framework.
- Additionally, in this tutorial I have combined some pre-processing and modelling steps into a single model pipeline.

```
# importing libraries necessary for data transformation, pre-processing, ml
model development etc

import pandas as pd
import numpy as np

from pyspark.sql.functions import desc, isnan, when, count, col, countDistinct,
regexp_replace
from pyspark.ml import Pipeline
from pyspark.ml.feature import Imputer, StringIndexer, OneHotEncoder,
VectorAssembler
from pyspark.ml.regression import GBTRRegressor
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
from pyspark.ml.evaluation import RegressionEvaluator
```

1. Data import

```
# Model development data is stored as Delta table stored in default database  
of my Databricks workspace.  
# Importing model development data as a spark dataframe, by using both SQL and  
Spark commands
```

```
df_sql = '''SELECT * FROM default.reg_data_missing'''  
df = spark.sql(df_sql)  
print('Dataframe data type : {}'.format(type(df)))  
print('Data Shape : {} Rows - {} Columns'.format((df.count()),  
(len(df.columns))))  
print('\nTop 5 rows')  
print(df.show(5))  
print('\nBottom 5 rows')  
print(df.orderBy(desc('ID')).show(5))  
print(df.show(5))
```

```
Dataframe data type : <class 'pyspark.sql.dataframe.DataFrame'>  
Data Shape : 6000 Rows - 9 Columns
```

```
Top 5 rows
```

ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
1	1.1	Ideal	H	SI1	VG	EX	GIA	5169.0
2	0.83	Ideal	H	VS1	ID	ID	AGSL	3470.0
3	0.85	Ideal	H	SI1	EX	EX	GIA	3183.0
4	0.91	Ideal	E	SI1	VG	VG	GIA	4370.0
5	0.83	Ideal	G	SI1	EX	EX	GIA	3171.0

```
only showing top 5 rows
```

```
None
```

```
Bottom 5 rows
```

ID	Carat Weight	Cut	Color	Clarity	Polish	Symmetry	Report	Price
1	1.1	Ideal	H	SI1	VG	EX	GIA	5169.0
2	0.83	Ideal	H	VS1	ID	ID	AGSL	3470.0
3	0.85	Ideal	H	SI1	EX	EX	GIA	3183.0
4	0.91	Ideal	E	SI1	VG	VG	GIA	4370.0
5	0.83	Ideal	G	SI1	EX	EX	GIA	3171.0

```
print('All columns in the dataset with their datatypes')  
print(df.dtypes)
```

```
All columns in the dataset with their datatypes  
[('ID', 'int'), ('Carat Weight', 'double'), ('Cut', 'string'), ('Color', 'stri
```

```

ng'), ('Clarity', 'string'), ('Polish', 'string'), ('Symmetry', 'string'), ('Report', 'string'), ('Price', 'double')]

# replacing all the whitespaces with '_' in all the column names

column_names = [column.replace(' ', '_') for column in df.columns]
df = df.toDF(*column_names)
print(df.show(2))

+---+-----+-----+-----+-----+-----+-----+
| ID|Carat_Weight| Cut|Color|Clarity|Polish|Symmetry|Report| Price|
+---+-----+-----+-----+-----+-----+-----+
|  1|      1.1|Ideal|    H|   SI1|     VG|      EX|    GIA| 5169.0|
|  2|      0.83|Ideal|    H|   VS1|     ID|      ID| AGSL| 3470.0|
+---+-----+-----+-----+-----+-----+-----+
only showing top 2 rows

```

None

2. Summary Statistics and Data Imputation on model development data

```

# collecting all the numerical and categorical variables and saving them in two different lists

categorical_variables = [item[0] for item in df.dtypes if
item[1].startswith('string')]
print('All categorical variables')
print(categorical_variables)
numerical_variables = [item[0] for item in df.dtypes if
item[1].startswith('double') or item[1].startswith('int')]
print('All numeric variables')
print(numerical_variables)

All categorical variables
['Cut', 'Color', 'Clarity', 'Polish', 'Symmetry', 'Report']
All numeric variables
['ID', 'Carat_Weight', 'Price']

```

```
# counting null/nan/missing values in all the columns

dict_1 = dict(df.dtypes)
df_info = pd.DataFrame(dict_1.items(), columns=['Column', 'DataType'])
null_count = df.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for
c in df.columns]).toPandas()
null_count = list(null_count.iloc[0,:])
df_info['NullCount'] = null_count
print(df_info)
```

	Column	DataType	NullCount
0	ID	int	0
1	Carat_Weight	double	30
2	Cut	string	0
3	Color	string	0
4	Clarity	string	0
5	Polish	string	0
6	Symmetry	string	30
7	Report	string	0
8	Price	double	0

```
# checking if any of the numerical columns has a missing values
```

```
missing_numerical_variables = list(df_info['Column'].loc[(df_info['NullCount'] > 0) & (df_info['DataType'] != 'string')])  
missing_numerical_variables
```

```
Out[7]: ['Carat_Weight']
```

```
# summary statistics on all the numerical columns
```

```
print(df.select([column for column in numerical_variables if column != 'ID']).describe().show())
```

summary	Carat_Weight	Price
count	5970	6000
mean	1.3346180904522886	11791.579333333333
stddev	0.4757117705068891	10184.350050741188
min	0.75	2184.0
max	2.91	101561.0

```
None
```

```
# summary statistics on all the numerical columns

print(df.agg(*([countDistinct(col(c)).alias(c) for c in
categorical_variables])).show())
```

```
+----+----+----+----+----+
|Cut|Color|Clarity|Polish|Symmetry|Report|
+----+----+----+----+----+
|  5|    6|     7|     4|      4|     2|
+----+----+----+----+----+
```

None

```
# counting the frequency of each sub-level for all the categorical columns :
null values identified
```

```
for column in categorical_variables:
    df.groupBy(column).count().orderBy('count').show()
```

```
+-----+----+
|          Cut|count|
+-----+----+
|      Fair|  129|
|Signature-Ideal| 253|
|      Good|  708|
|  Very Good| 2428|
|      Ideal| 2482|
+-----+----+
```

```
+----+----+
|Color|count|
+----+----+
|    D|  661|
|    E|  778|
|    I|  968|
|    F| 1013|
|    H| 1079|
|    G| 1501|
+----+----+
```

```

# replacing whitespace and hyphen characters in all the sub-level values for
all the categorical variables

for column in categorical_variables:
    df = df.withColumn('{{}'.format(column), regexp_replace('{{}'.format(column),
    '[\s-]', '_')}

for column in categorical_variables:
    df.select(column).distinct().show()

```

```

+-----+
|      Cut|
+-----+
|  Very_Good|
|Signature_Ideal|
|      Ideal|
|      Good|
|      Fair|
+-----+

+----+
|Color|
+----+
|   F|
|   E|
|   D|
|   G|
|   I|
|   H|
+----+

```

```
# checking all the rows where Carat_Weight value is missing
```

```
print(df.filter(col('Carat_Weight').isNull()).show())
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
--+
| ID|Carat_Weight|          Cut|Color|Clarity|Polish|Symmetry|Report|  Pric
e|
+-----+-----+-----+-----+-----+-----+-----+-----+
--+
| 121|        null|      Good|    H|    SI1|     EX|     EX|    GIA| 4563.
0|
| 405|        null|      Good|    I|    SI1|     EX|      G|    GIA|15743.
0|

```

591	null	Fair	D	SI1	EX	EX	GIA	4537.
0								
640	null Signature_Ideal		I	SI1	ID	ID	AGSL	11110.
0								
816	null	Very_Good	H	VS2	EX	VG	GIA	5571.
0								
878	null	Very_Good	I	SI1	G	null	GIA	4664.
0								
1162	null	Very_Good	I	VS2	VG	VG	GIA	15895.
0								

```
# imputing the missing Carat_Weight rows using PySpark's Imputer functions
df_impute = df
imputer = Imputer(inputCols=missing_numerical_variables, outputCols=
['{}_imputed'.format(c) for c in missing_numerical_variables])
df_impute = imputer.fit(df).transform(df_impute)

# checking all the rows where Carat_Weight value is missing : all missing
values imputed
print(df_impute.show(5))

+---+-----+---+---+---+---+---+---+---+
| ID|Carat_Weight| Cut|Color|Clarity|Polish|Symmetry|Report| Price|Carat_Weig
ht_imputed|
+---+-----+---+---+---+---+---+---+---+
|  1|      1.1|Ideal| H| SI1| VG| EX| GIA|5169.0|
1.1|
|  2|      0.83|Ideal| H| VS1| ID| ID| AGSL|3470.0|
0.83|
|  3|      0.85|Ideal| H| SI1| EX| EX| GIA|3183.0|
0.85|
|  4|      0.91|Ideal| E| SI1| VG| VG| GIA|4370.0|
0.91|
|  5|      0.83|Ideal| G| SI1| EX| EX| GIA|3171.0|
0.83|
+---+-----+---+---+---+---+---+---+---+
-----+
only showing top 5 rows
```

None

```

# checking all the rows where Symmetry value is missing
print(df_impute.filter(col('Carat_Weight').isNull()).show(5))

# dropping the column Carat_Weight as all the original and imputed values are
# saved in Carat_Weight_imputed column
df_impute = df_impute.drop('Carat_Weight')

+-----+-----+-----+-----+-----+-----+-----+
+-----+
| ID|Carat_Weight|          Cut|Color|Clarity|Polish|Symmetry|Report|  Price
|Carat_Weight_imputed|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
| 121|      null|      Good|    H|    SI1|    EX|    EX|    GIA| 4563.0
| 1.3346180904522886|
| 405|      null|      Good|    I|    SI1|    EX|     G|    GIA|15743.0
| 1.3346180904522886|
| 591|      null|      Fair|    D|    SI1|    EX|    EX|    GIA| 4537.0
| 1.3346180904522886|
| 640|      null|Signature_Ideal|    I|    SI1|    ID|    ID|    AGSL|11110.0
| 1.3346180904522886|
| 816|      null|Very_Good|    H|    VS2|    EX|    VG|    GIA| 5571.0
| 1.3346180904522886|
+-----+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows

```

None

```
print(df_impute.filter(col('Symmetry').isNull()).show())
```

ID	Cut	Color	Clarity	Polish	Symmetry	Report	Price	Carat_Weight_imputed
71	Ideal	G	VS2	EX	null	GIA	6264.0	1.02
141	Ideal	I	SI1	EX	null	GIA	16718.0	2.04
191	Very_Good	G	VVS2	G	null	GIA	8013.0	1.02
244	Ideal	F	IF	EX	null	GIA	11327.0	1.07
309	Ideal	H	SI1	VG	null	GIA	5529.0	

```

1.15|
| 878|Very_Good|     I|      SI1|      G|      null|    GIA| 4664.0|  1.3346180904522
886|
| 975|     Good|     F|     VVS2|     EX|      null|    GIA|17203.0|
1.51|

```

```
# imputing all missing Symmetry values with mode Symmetry value : VG in this case
```

```

temp_df= df_impute.groupBy('Symmetry').count()
mode_variable = temp_df.orderBy(temp_df['count'].desc()).collect()[0][0]
mode_count_value = temp_df.orderBy(temp_df['count'].desc()).collect()[0][1]
print('Mode Value : {} \n'.format(mode_variable), 'Mode Count : {} \n'.format(mode_count_value))
df_impute = df_impute.fillna({'Symmetry':mode_variable})
df_impute = df_impute.withColumnRenamed('Symmetry', 'Symmetry_imputed')
print(df_impute.show(5))

```

```

Mode Value : VG
Mode Count : 2401
+---+-----+-----+-----+-----+-----+-----+-----+
----+
| ID| Cut|Color|Clarity|Polish|Symmetry_imputed|Report| Price|Carat_Weight_imputed|
+---+-----+-----+-----+-----+-----+-----+-----+
----+
|  1|Ideal|   H|    SI1|     VG|                  EX|    GIA|5169.0|
1.1|
|  2|Ideal|   H|    VS1|     ID|                  ID|    AGSL|3470.0|
0.83|
|  3|Ideal|   H|    SI1|     EX|                  EX|    GIA|3183.0|
0.85|
|  4|Ideal|   E|    SI1|     VG|                  VG|    GIA|4370.0|
0.91|
|  5|Ideal|   G|    SI1|     EX|                  EX|    GIA|3171.0|
0.83|
+---+-----+-----+-----+-----+-----+-----+-----+
----+
only showing top 5 rows

```

3. Visual analysis using some SQL and Databricks inbuilt plot functions

```

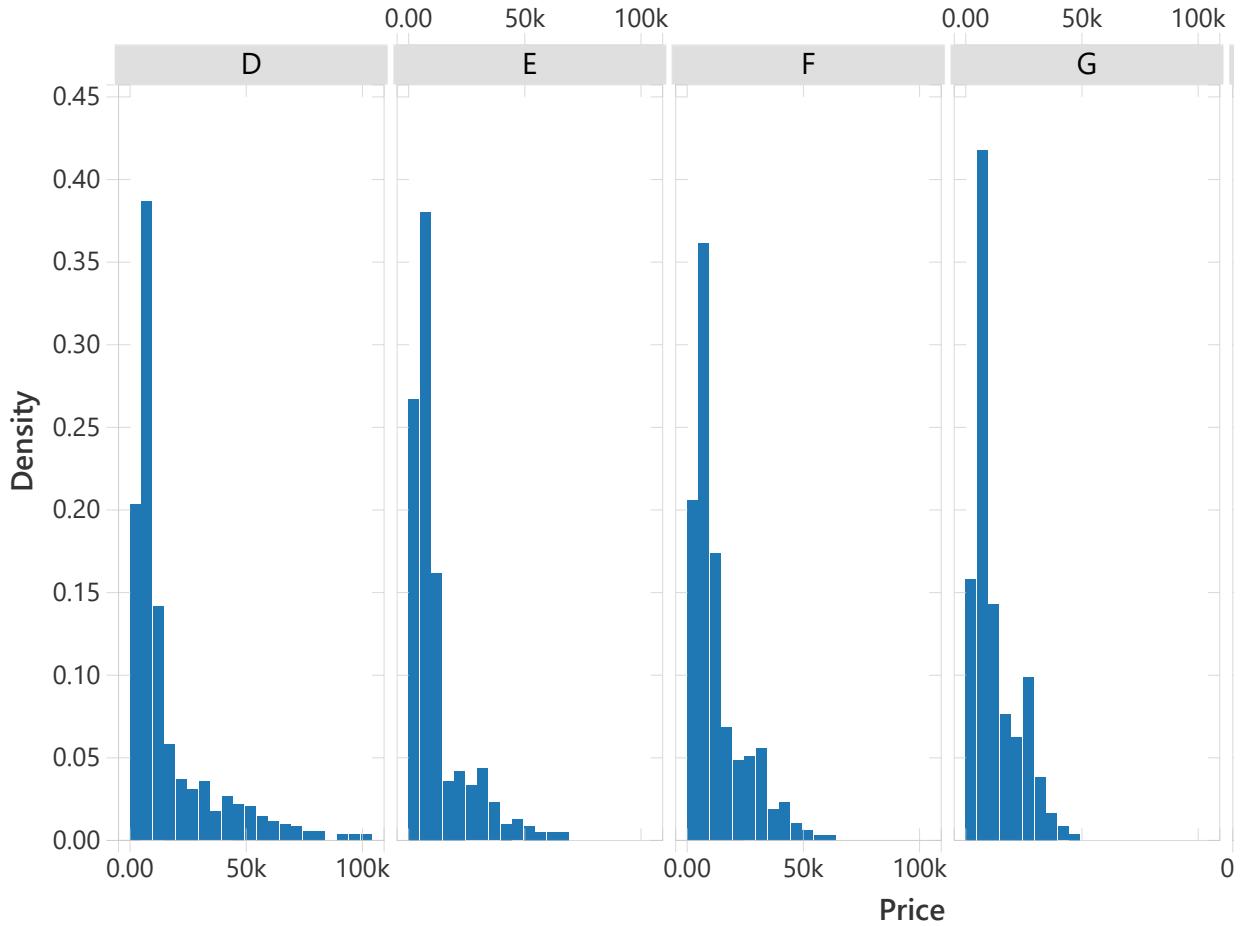
# creating separate lists for categorical and numerical variables

categorical_variables = [item[0] for item in df_impute.dtypes if
item[1].startswith('string')]

numerical_variables = [item[0] for item in df_impute.dtypes if
item[1].startswith('double') or item[1].startswith('int')]

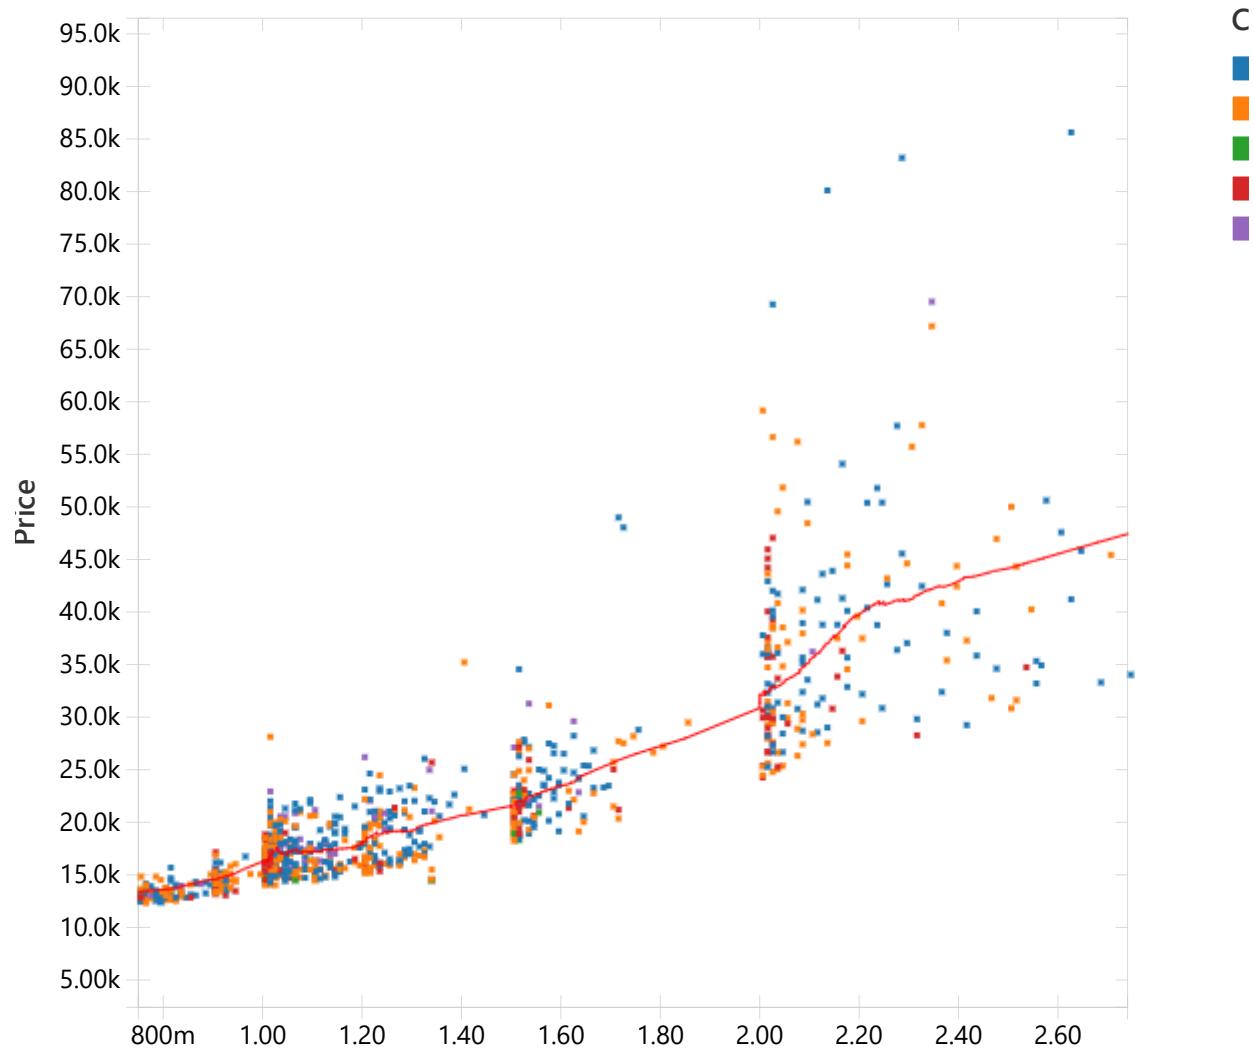
%sql
SELECT ID, Price, Cut, Color
FROM default.reg_data_missing;

```



Aggregated (by count) in the backend

```
display(df_impute)
```



4. Data Pre-processing and Model development

```
# Splitting model development data into train-test data  
  
df_train, df_test = df_impute.randomSplit([0.8, 0.2], seed=21)  
(df_train.cache().count(), df_test.count())  
  
Out[19]: (4746, 1254)
```

```

# creating StringIndexer and OneHotEncoder objects to convert encode
categorical data
# creating VectorAssembler objects to collect all the encoded and numerical
features into a sparse vector
# saving all the objects into a 'stage' list, which will be passed into a model
developement pipeline

stages = []

for column in categorical_variables:
    string_indexer = StringIndexer(inputCol = column, outputCol =
column+'_Index')
    oh_encoder = OneHotEncoder(inputCols=[string_indexer.getOutputCol()],
outputCols=[column+'_classVec'], dropLast=True)
    stages +=[string_indexer, oh_encoder]

assemble_inputs = [c+'_classVec' for c in categorical_variables] +
numerical_variables
assemble_inputs.remove('ID')
vectorAssembler = VectorAssembler(inputCols=assemble_inputs,
outputCol='x_features')
stages +=[vectorAssembler]

# creating an object for a base Gradient Boosting regressor model
gbt = GBTRRegressor(labelCol='Price', featuresCol='x_features', maxIter=15,
maxDepth=5, seed=21)

# adding model object to 'stages' list
stages.append(gbt)

# passing the 'stages' to a Pipeline object
pipeline = Pipeline(stages=stages)
print('Pipeline with stages')
print(pipeline)
print(stages)

Pipeline with stages
Pipeline_9ad1332d7939
[StringIndexer_adf2904cb17a, OneHotEncoder_ea22da22ef55, StringIndexer_adc6af6
25fee, OneHotEncoder_76738f3743c0, StringIndexer_84c0d334203a, OneHotEncoder_0
310e5531a97, StringIndexer_8b9906fc1c84, OneHotEncoder_425ac27bc92e, StringInd
exer_917bf2c8a455, OneHotEncoder_97fdd39e2fd0, StringIndexer_b092586198d4, One
HotEncoder_cbd148b27d36, VectorAssembler_5dd87a1b3005, GBTRRegressor_9ff002c9b0
ca]

```

```

# fitting the Pipeline object on train data : all the pre-processing and and
model fitting will occur in this step
pipelineModel = pipeline.fit(df_train)

# fitting the trained pre-processing and model-fit pipeline on test data to
calculate predictions on test data
gbt_predictions = pipelineModel.transform(df_test)

# selecting variables to check the predictions on test data
gbt_predictions.select('ID', 'x_features', 'Price', 'prediction').take(3)

Out[22]: [Row(ID=12, x_features=SparseVector(24, {2: 1.0, 8: 1.0, 9: 1.0, 20:
1.0, 21: 1.0, 22: 5161.0, 23: 1.01}), Price=5161.0, prediction=5306.9535728705
9),
Row(ID=22, x_features=SparseVector(24, {1: 1.0, 6: 1.0, 11: 1.0, 15: 1.0, 20:
1.0, 21: 1.0, 22: 3559.0, 23: 0.8}), Price=3559.0, prediction=3813.52864952471
24),
Row(ID=41, x_features=SparseVector(24, {2: 1.0, 4: 1.0, 9: 1.0, 16: 1.0, 20:
1.0, 21: 1.0, 22: 4905.0, 23: 1.01}), Price=4905.0, prediction=4566.6853523534
1)]]

# defining MAE as metric to evaluate model performance
# evaluating model predictions on test-data using Mean Absolute Error metric

gbt_evaluator = RegressionEvaluator(metricName='mae',
labelCol=gbt.getLabelCol(), predictionCol=gbt.getPredictionCol())
mae = round((gbt_evaluator.evaluate(gbt_predictions)),2)
print('Test MAE = {}'.format(mae))

Test MAE = 594.62

```

```

# defining model hyperparameter grid
parameter_grid = (ParamGridBuilder()
    .addGrid(gbt.maxDepth, [4,5,6])
    .addGrid(gbt.maxIter, [10, 15, 20])
    .build())

# defining cross-validation object with model, hyperparameter grid, and evaluator
# object values
cv = CrossValidator(estimator=pipeline, evaluator=gbt_evaluator,
estimatorParamMaps=parameter_grid, numFolds=3, parallelism = 4)

# fitting model on train data and calculating predictions on test data
gbt_cv_model = cv.fit(df_train)
gbt_cv_predictions = gbt_cv_model.transform(df_test)
gbt_cv_predictions.select('ID', 'x_features', 'Price', 'prediction').take(3)

Out[24]: [Row(ID=12, x_features=SparseVector(24, {2: 1.0, 8: 1.0, 9: 1.0, 20: 1.0, 21: 1.0, 22: 5161.0, 23: 1.01}), Price=5161.0, prediction=5068.856103818482),
Row(ID=22, x_features=SparseVector(24, {1: 1.0, 6: 1.0, 11: 1.0, 15: 1.0, 20: 1.0, 21: 1.0, 22: 3559.0, 23: 0.8}), Price=3559.0, prediction=3732.5177310762324),
Row(ID=41, x_features=SparseVector(24, {2: 1.0, 4: 1.0, 9: 1.0, 16: 1.0, 20: 1.0, 21: 1.0, 22: 4905.0, 23: 1.01}), Price=4905.0, prediction=4687.2772597766725)]

```

Test MAE = 525.19